

By Jiahui Tang

Reported Optimized Result:

- (1) GCC -Ofast:
 - (1) Best elapsed time: 2.50645 s
 - (2) Best real time: 0.195s
- (2) ICC -fast:
 - (1) Best elapsed time: 0.75382 s
 - (2) Best real time: 0.129s

Best improvements:

- (1) in terms of elapsed time: 24.606x faster
- (2) In terms of real time: 144.193x faster

No optimization:

```
collect2: error: ld returned 1 exit status
ubuntu@ip-172-31-77-30:~$ gcc -DUSE_CLOCK slowpoke.c -o slowpoke -lm
ubuntu@ip-172-31-77-30:~$ time ./slowpoke
*****
||c||_F = 2630354827429312.50
Elapsed Time: 18.5493 s.

real    0m18.601s
user    0m18.569s
sys     0m0.032s
ubuntu@ip-172-31-77-30:~$
```

--

Details:

Version 1 - GCC with slightly lower result marginal error but executed slower

1. Launch m5.2xlarge instance
 2. To ensure that it doesn't attempt to sync j and k between threads, we could set them to private versions. So I also moved declaration into the loop body to parallelize them.
 3. Change matrix b's index from b[k][j] to b[j][k] to improve memory locality;
- Since matrix b[i][j]= I*j thus it is a symmetric matrix, so we can directly switch j and k;
3. 4* unrolling the inner loop

```
get_time(&tstart);
#pragma omp parallel for
for (i=0; i<N; i++)
{
    for(int jj=0; jj<N; jj++)
        for (int kk=0; kk<N; kk+=4) {
            c[i][jj] += a[i][kk] * b[jj][kk] + a[i][kk+1] * b[jj][kk+1] + a[i][kk+2] * b[jj][kk+2] + a[i][kk+3] * b[jj][kk+3];
        }
}
get_time(&tend);
```

4. add OpenMP tags as indicated #pragma omp parallel for

3. set thread = 16;

3. Compile using

```
##gcc -fopenmp -Ofast -march=native -DCLOCK_REALTIME seq_mm.c timing.c -o omp_seq
```

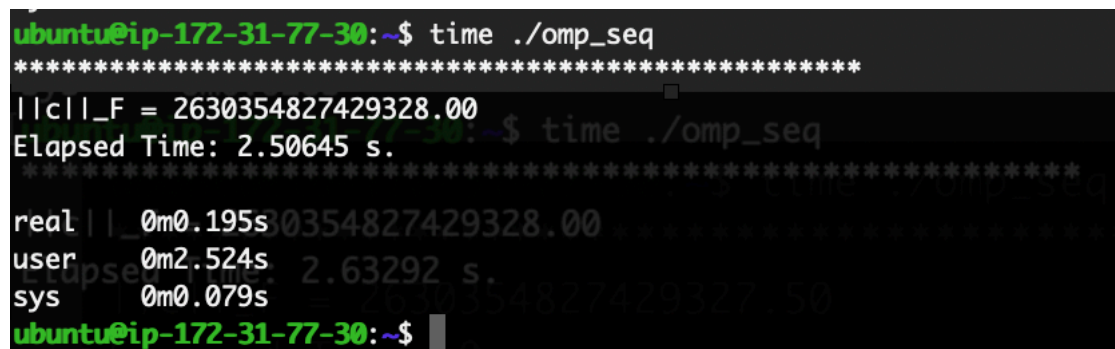
```
gcc -fopenmp -Ofast -march=native -DUSE_CLOCK seq_mm.c -o omp_seq
```

Result

```
ubuntu@ip-172-31-77-30:~$ time ./omp_seq
*****
||c||_F = 2630354827429328.00
Elapsed Time: 2.50645 s.
```

```
real 0m0.195s
user 0m2.524s
sys 0m0.079s
```

real time reported with screenshot:

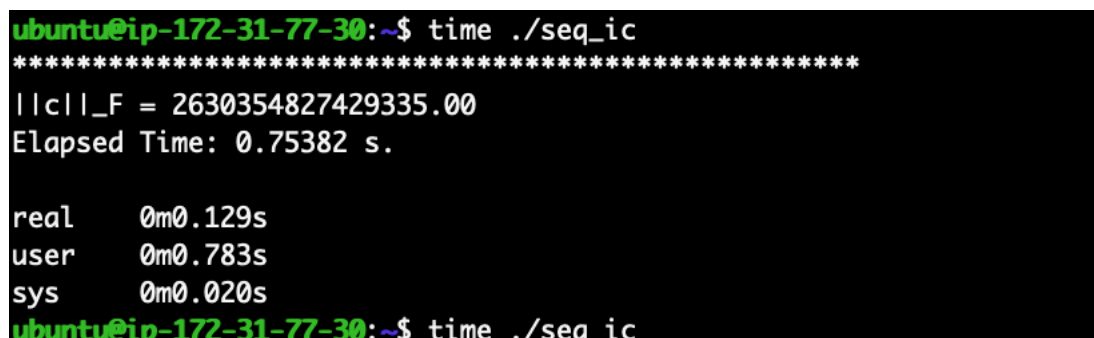


```
ubuntu@ip-172-31-77-30:~$ time ./omp_seq
*****
||c||_F = 2630354827429328.00
Elapsed Time: 2.50645 s.
*****
real    0m0.195s
user    0m2.524s
sys     0m0.079s
ubuntu@ip-172-31-77-30:~$
```

--

Version 2 (ICC- faster but with slightly more result marginal error)

1. Add apt repo via <https://software.intel.com/content/www/us/en/develop/articles/installing-intel-oneapi-toolkits-via-apt.html>
2. Run
apt-get install intel-oneapi-compiler-dpcpp-cpp-and-cpp-classic
3. Run
source /opt/intel/oneapi/setvars.sh
4. Compile by:
icc -fopenmp -fast -DUSE_CLOCK ./seq_mm.c -o seq_ic
5. Screenshots:



```
ubuntu@ip-172-31-77-30:~$ time ./seq_ic
*****
||c||_F = 2630354827429335.00
Elapsed Time: 0.75382 s.
*****
real    0m0.129s
user    0m0.783s
sys     0m0.020s
ubuntu@ip-172-31-77-30:~$ time ./seq_ic
```

```

ubuntu@ip-172-31-77-30:~$ OMP_NUM_THREADS=16
ubuntu@ip-172-31-77-30:~$ time ./seq_ic
*****
||c||_F = 2630354827429335.00
Elapsed Time: 1.48407 s.

real    0m0.129s
user    0m1.512s
sys     0m0.036s
ubuntu@ip-172-31-77-30:~$ █

```

Attachements:

8 threads:

```
ubuntu@ip-172-31-77-30:~$ OMP_NUM_THREADS=8 time ./seq_ic
```

```
*****
```

```
||c||_F = 2630354827429335.00
```

```
Elapsed Time: 0.962459 s.
```

```
0.98user 0.02system 0:00.15elapsed 649%CPU (0avgtext+0avgdata
55504maxresident)k
```

```
0inputs+0outputs (0major+13461minor)pagefaults 0swaps
```

16 threads:

```
ubuntu@ip-172-31-77-30:~$ OMP_NUM_THREADS=16 time ./seq_ic
```

```
*****
```

```
||c||_F = 2630354827429335.00
```



```
Elapsed Time: 1.48194 s.
```

```
1.52user 0.02system 0:00.12elapsed 1207%CPU (0avgtext+0avgdata
55856maxresident)k
```

```
0inputs+0outputs (0major+13549minor)pagefaults 0swaps
```

Meme:







4 X unrolling!



Try more unrolling!



Oops, No, back to
4 X unrolling!



Turn around, switch j
and k index!



Meme Reference:

Meawbin

<http://meawbinneko.com/meawbincomic.html>